

Execution Concurrency in Python

Concurrency in Python

- Two kinds of tasks that can be used for concurrency:
 - Processes
 - Python multiprocessing module
 - have their own memory (e.g., variables)
 - A process has one or more threads
 - Threads
 - Python threading module
 - share the memory of the process that they are part of
 - are considered more lightweight in that they can spawn very quickly

Multiprocessing Pros/Cons

- Pros:
 - Separate memory space
 - Code is generally more straightforward
 - Takes better advantage of multiple CPUs/cores
 - Child processes can be interrupted/killed
- Cons:
 - Larger memory footprint
 - Interprocess communications is more complicated, i.e., IPC versus shared memory objects.

Threading Pros/Cons

- Pros:
 - Lightweight, i.e., low memory footprint
 - Shared memory that makes access to state from another thread easier
- Cons:
 - Code is a bit more complicated due to memory sharing.

Python Threading

Python Threads

- Enables multiple execution threads so that different tasks can be performed (almost) concurrently, i.e., in parallel (it is really timesharing the processor between threads).
- Example: a python network server can create a thread whenever a new client connects. This helps avoid some of the socket blocking behaviour that might otherwise occur.

Python Threading Module

- Import the threading module:

```
import threading
```

- Create a new thread:

```
t = threading.Thread(target=handler, args=(i,))
```

where handler is the function to be threaded and args is a tuple of arguments passed to the function (Note that the comma may be needed in **(i,)** above to ensure that args is a tuple). Other arguments include naming the thread, e.g., **name=<name>**.

- Normally one keeps a list of the created threads:

```
thread_list = [ ]
```

Python Threading Module

- Add the new thread to the thread list:

`thread_list.append(new_thread)`

- We must start the execution of the thread, once it is created.

`new_thread.start()`

- You can wait until a thread terminates. This is a way of synchronizing your code, e.g.,

`new_thread.join()`

This blocks the calling thread until the called thread terminates.

Python Threading Module

- Daemon vs non-daemon threads:

`new_thread.daemon = True # or False`

- When true, it means that the thread will be terminated when the calling script exits. Otherwise the script will not exit until any threads have completed execution.
- When true, the calling script doesn't have to worry about terminating threads that it has created.

Python Threading Examples

- See `threading_example.py`.
- `daemon_thread_illustration.py`
- See `ping_threading_example.py`
- See `EchoClientServer_Thread.py`

Synchronizing Access

- When using threads it is important to avoid conflicts when multiple threads need access to a single variable or resource.
- Overlapping accesses/modifications may cause all kinds of problems that can be load dependent and difficult to debug.
- One way of handling this is using atomic operations. In Python the follow are atomic:
 - reading/replacing a single instance attribute
 - reading/replacing a single global variable
 - getting an item from a list
 - modifying a list in place (e.g. using "append")
 - fetching an item from a dictionary
 - modifying a dictionary in place (e.g. adding an item, or calling the clear method)

Mutual Exclusion via Locking

- Note that operations that read a variable, modifies it, then writes it back are not atomic!
- It is safest to use locking, which is part of the threading module. This works as follows:

```
lock = threading.Lock()
```

```
...
```

```
lock.acquire() # this operation will block if  
               # the lock is already held
```

```
... The shared resource is locked. Do whatever  
you want with it ...
```

```
lock.release() # release the lock so another  
               # thread can use the resource.
```

Locking

- In the previous example you need to make sure that the resource is released if something goes wrong. An alternate syntax:

```
lock = threading.Lock()
```

```
...
```

```
with lock:
```

```
    The shared resource is locked. Do whatever  
    you want with it. When you exit this block,  
    the lock is released.
```

- The advantage of this syntax is that `acquire()` and `release()` are automatically called when entering and leaving the "with" block. If something goes wrong, you will not be accidentally leaving the resource locked forever.

Locking

- `counter_function_variable.py`
- `counter_function_list.py`
- `counter_dict_locking_example.py`

Multiprocessing

Multiprocessing

- The basic multiprocessing module uses a similar API as the threading module. e.g.,

```
import multiprocessing
```

```
new_process = multiprocessing.Process(  
    target=self.connection_handler,  
    args=(connection,))
```

```
new_process.start()
```


Multiprocessing

- In Windows you need to include at the start of your Python 3 script:

```
if sys.platform == 'win32':  
    import multiprocessing.reduction
```

- You need to also call:

```
multiprocessing.freeze_support()
```

- See **EchoClientServer_Multiprocessing.py** example.

Multiprocessing

- This module contains a lot of sophisticated functionality, e.g.,
 - Queues and pipes for exchanging data between processes.
 - Contains equivalent of locking in thread module.
 - State can be shared between processes using shared memory